

# Kubernetes Basics

---

Getting Started With Google Kubernetes Engine

Version 1.5



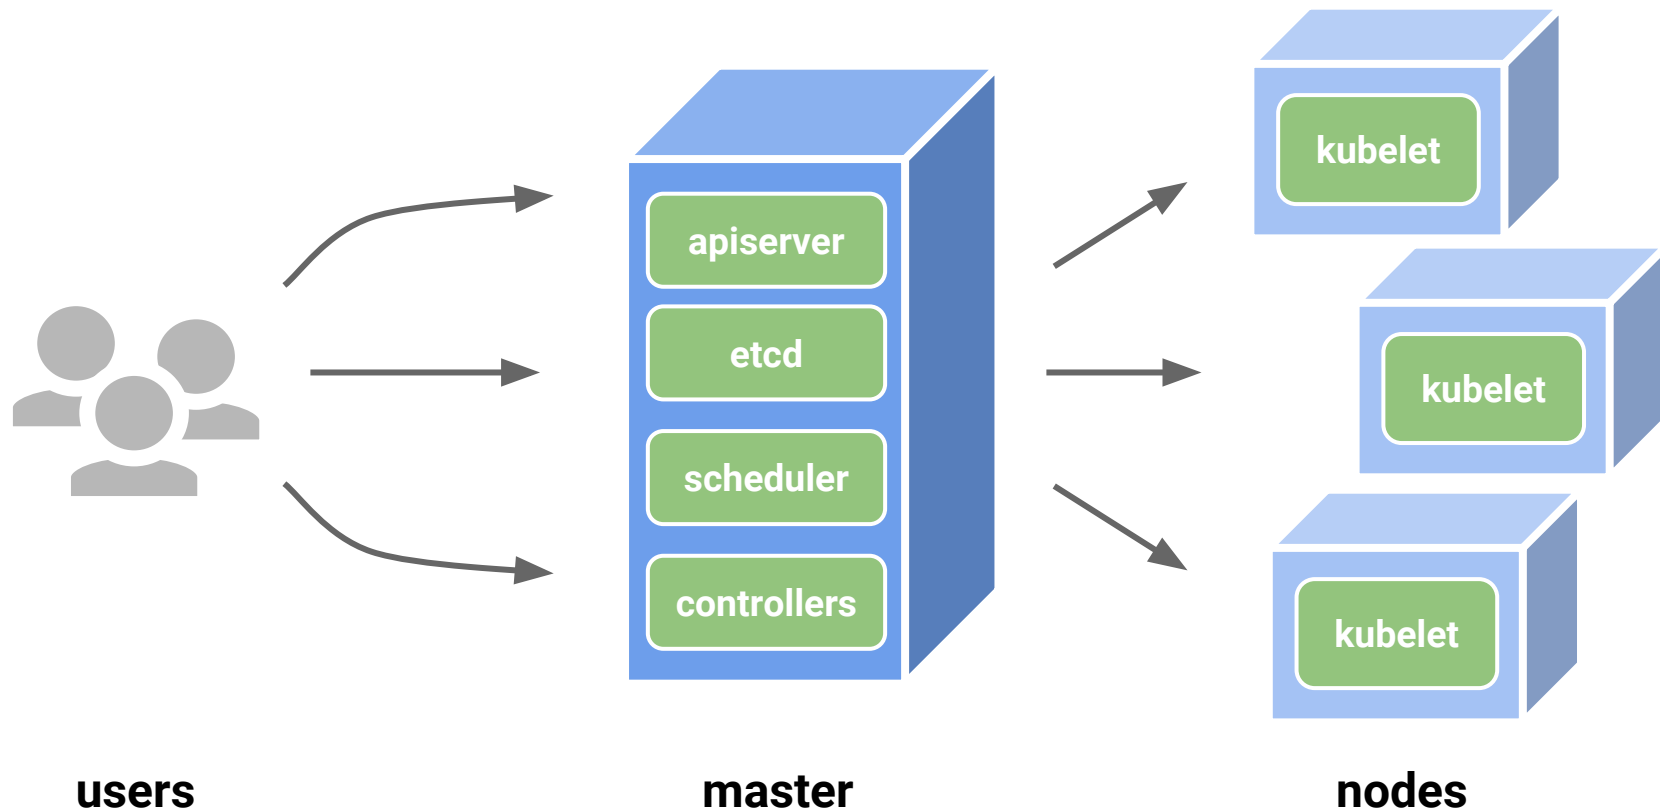
# Agenda

Clusters, nodes, and pods

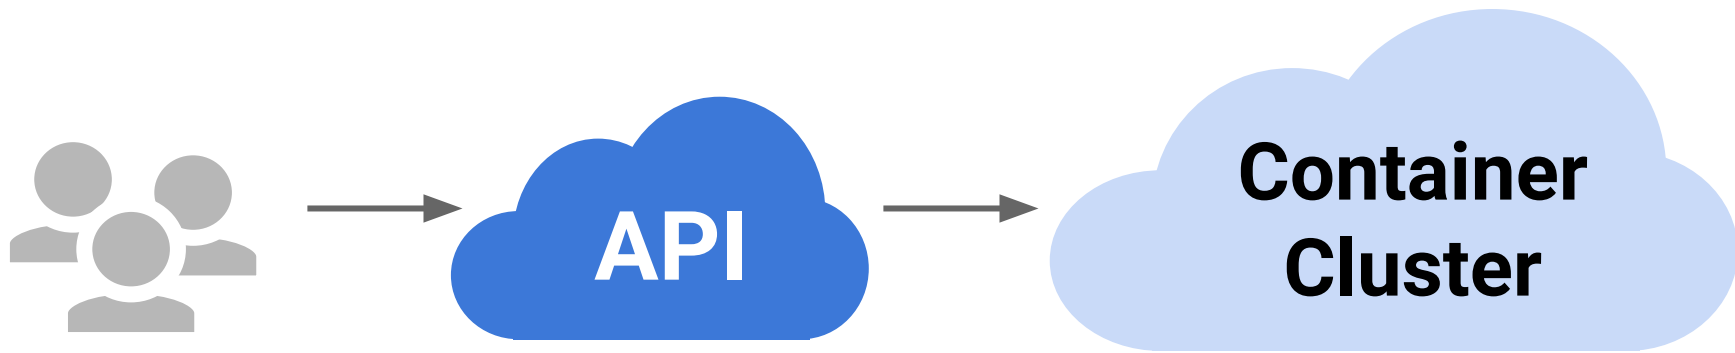
Services, labels and selectors

Volumes

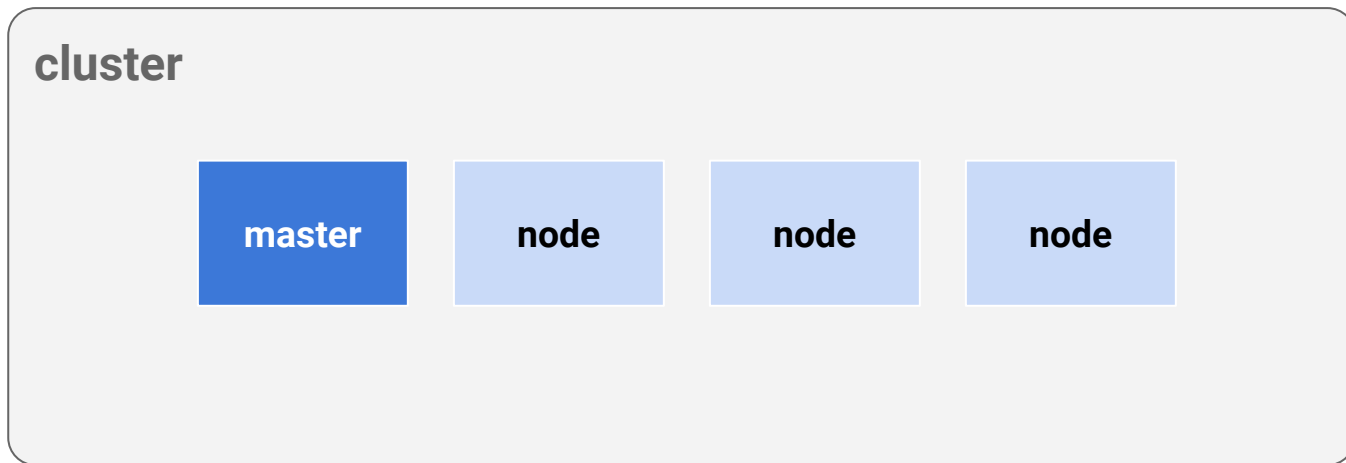
# The 10,000-foot view



# All users really care about



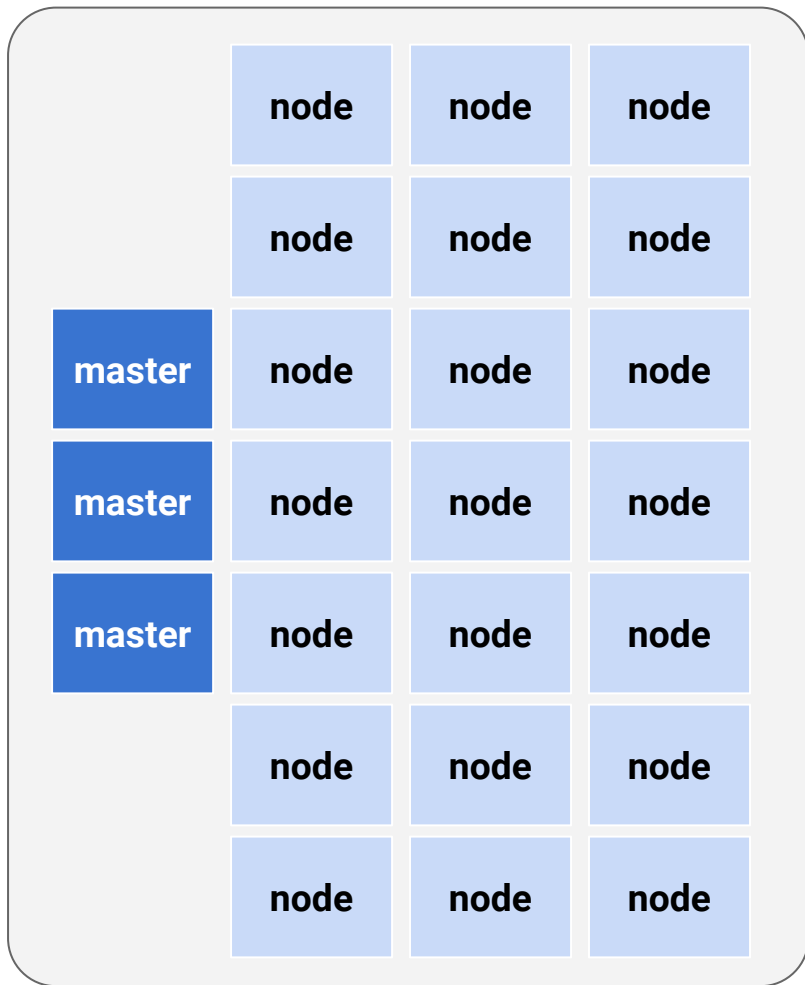
# A cluster is a set of computing instances that Kubernetes manages



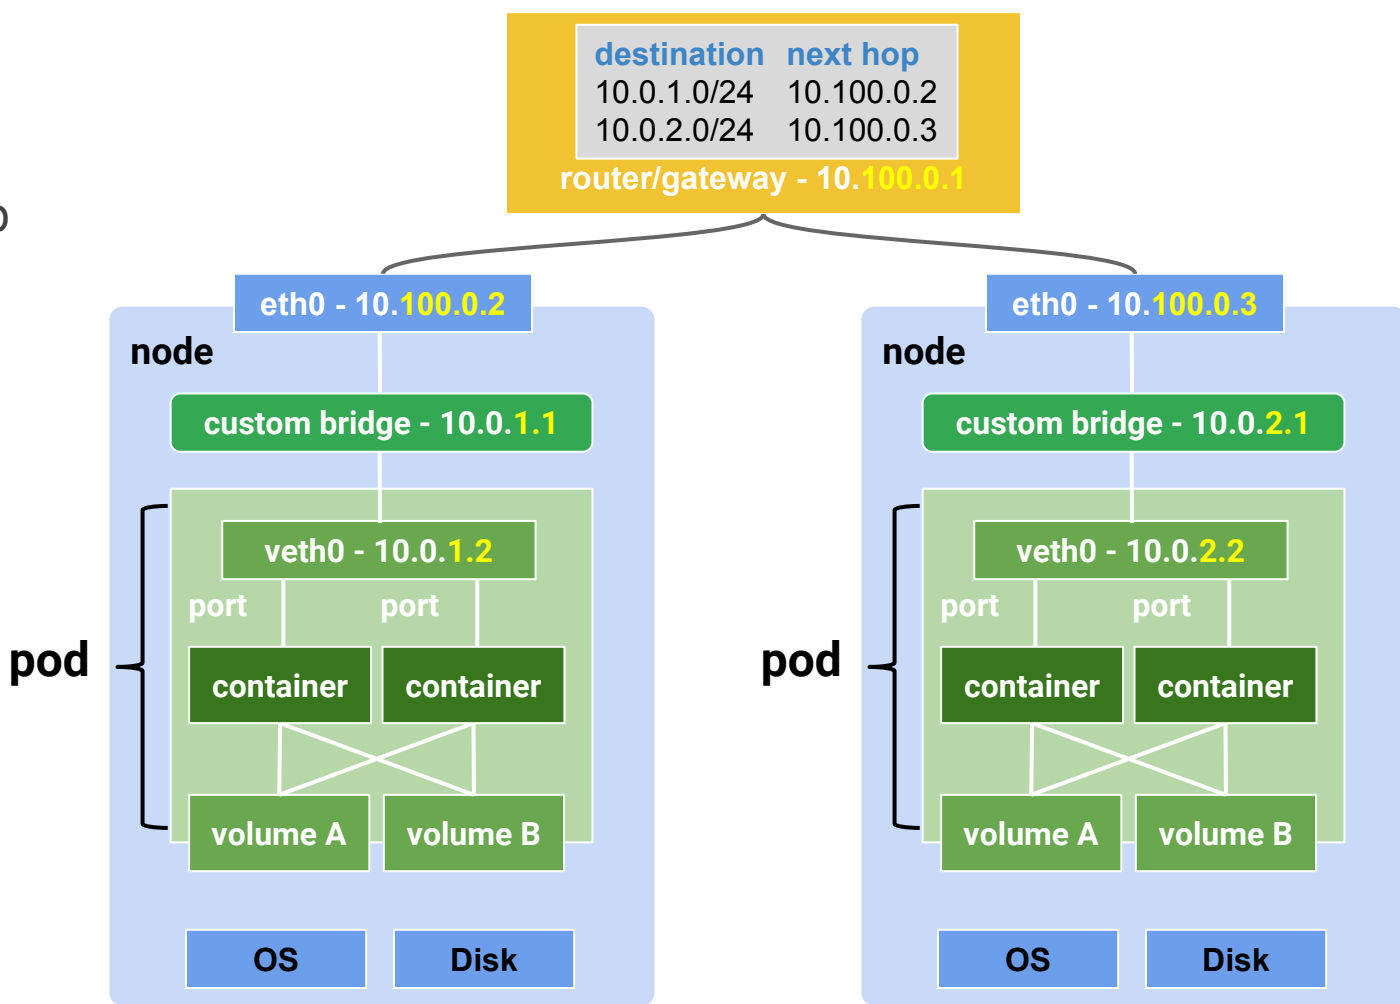
A cluster can have multiple masters and lots of nodes.

With regional clusters, masters and nodes are spread across three zones in a region for high availability.

By default, three masters are created per zone, but you can control the number.



A pod is  
analogous to a  
VM with a group  
of containers  
sharing  
networking and  
storage that are  
separate from  
the node

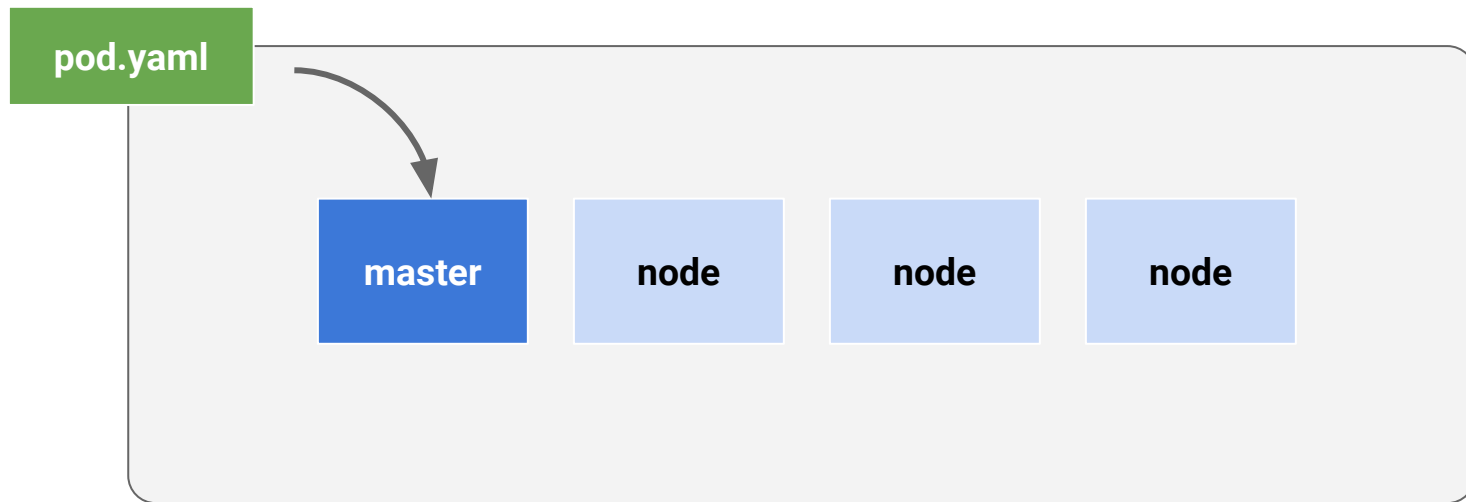


# You define a pod with a YAML file

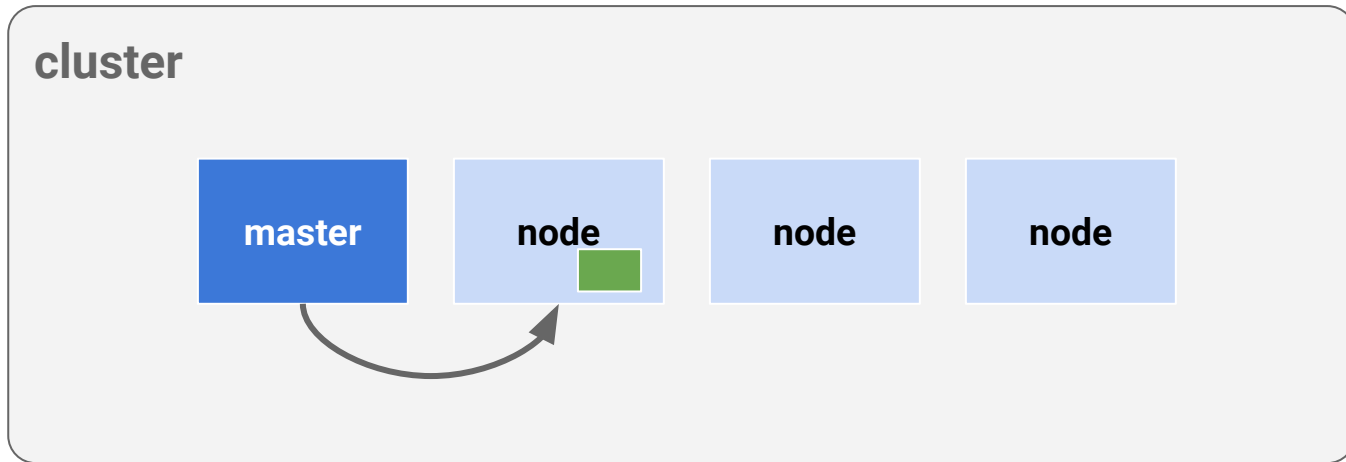
```
apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  containers:
    - name: my-app
      image: my-app
    - name: nginx-ssl
      image: nginx
  ports:
    - containerPort: 80
    - containerPort: 443
```



# You upload the YAML file to the master



And the master creates a pod on your set of nodes



# A pod file is composed of several parts; for example...

```
apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  containers:
    - name: my-app
      image: my-app
    - name: nginx-ssl
      image: nginx
  ports:
    - containerPort: 80
    - containerPort: 443
```

API version

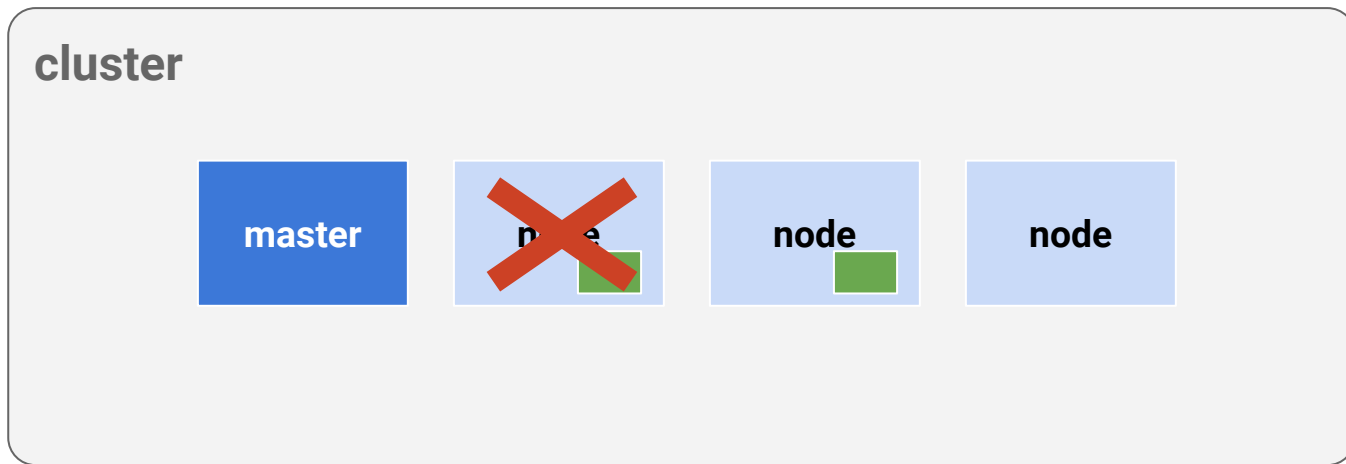
pod resource

pod name

two containers

NGINX front end on two ports

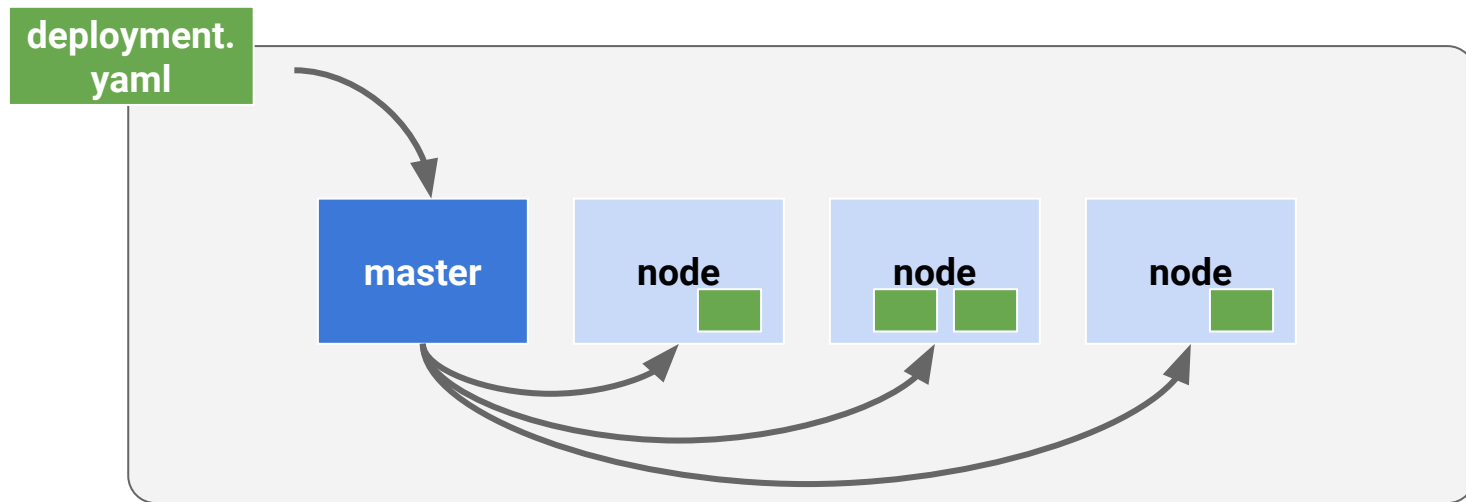
A deployment ensures that  $N$  pods are running in a cluster at any given time



```
kind: Deployment — deployment resource
apiVersion: v1.1
metadata:
  name: frontend — deployment name
spec:
  replicas: 4 — replicas
  selector: — pod selector
    role: web — role=web
  template:
    metadata:
      name: web
      labels: — pod label
        role: web — role=web
    spec:
      containers: — containers
      - name: my-app
        image: my-app
      - name: nginx-ssl
        image: nginx
        ports:
          - containerPort: 80
          - containerPort: 443
```

You define a deployment  
with a YAML file

You upload the YAML file to the master, and the scheduler decides where to run the pods



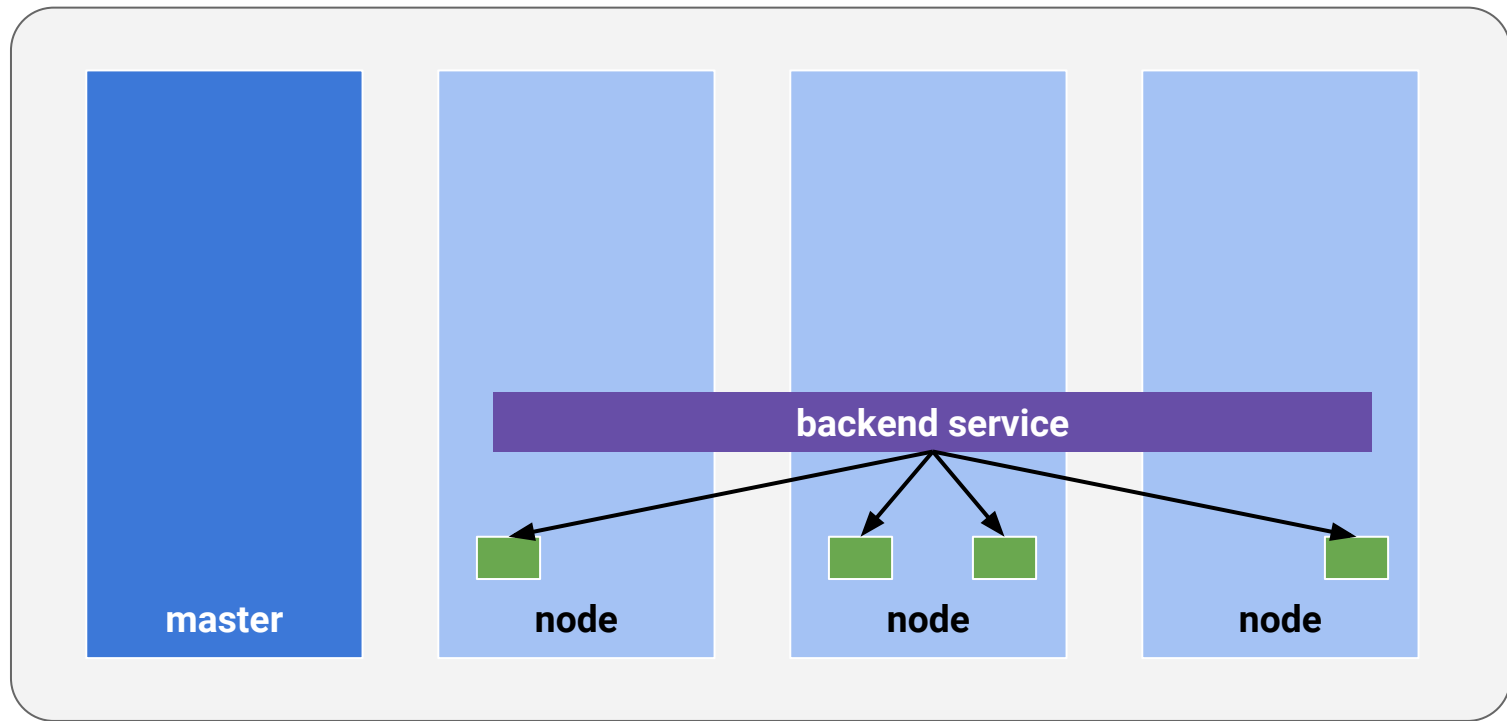
# Agenda

Clusters, nodes, and pods

Services, labels, and selectors

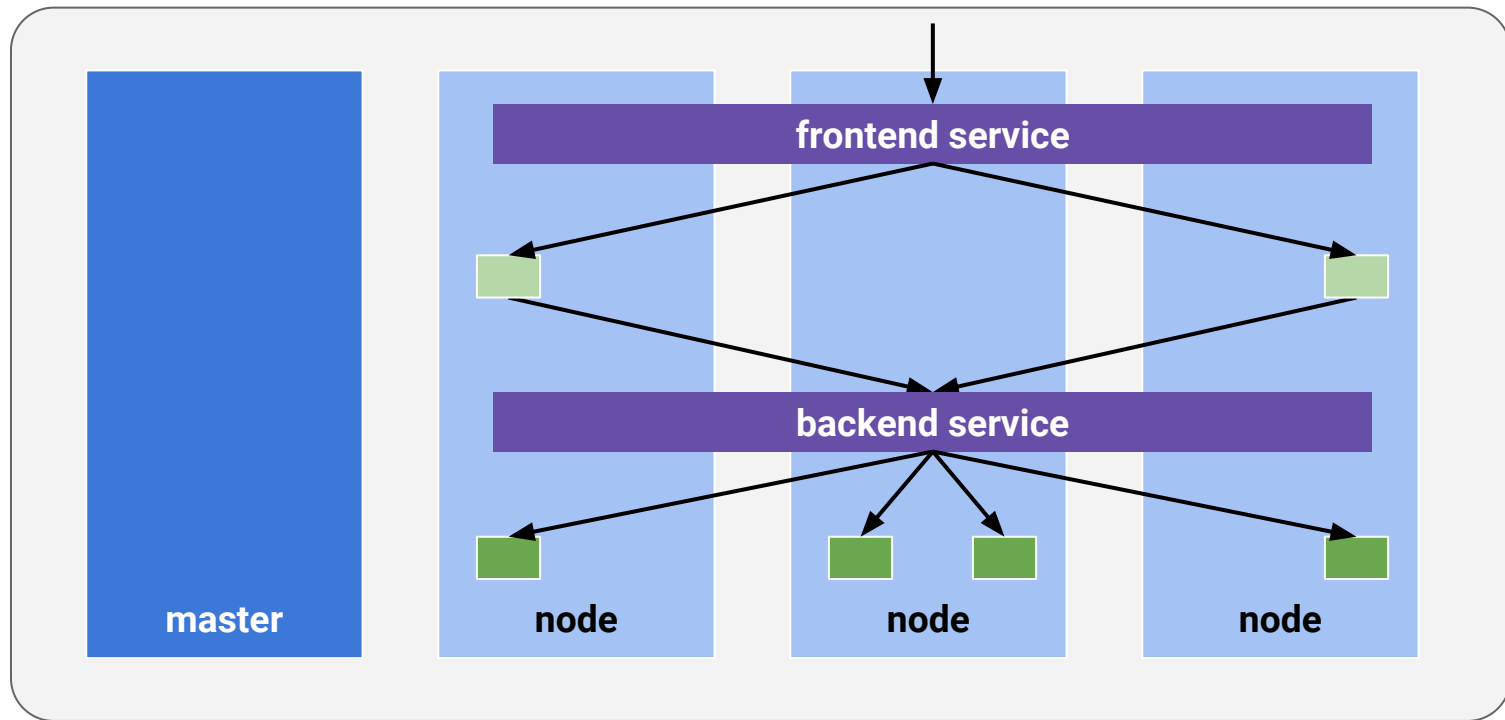
Volumes

A service assigns a fixed IP to your pod replicas and allows other pods or services to communicate with them





You can have multiple services with different configurations and features



# You define a service with a YAML file

```
kind: Service      — resource
apiVersion: v1
metadata:
  name: web-frontend
spec:
  ports:           — ports
  - name: http
    port: 80
    targetPort: 80
    protocol: TCP
  selector:        — pod selector
    role: web
  type: LoadBalancer — type is load balancer
```

# Labels are metadata you can assign to any API object and represent identity

They are

- The only grouping mechanism for pods
- Search by selectors



# This example has four pods and three labels

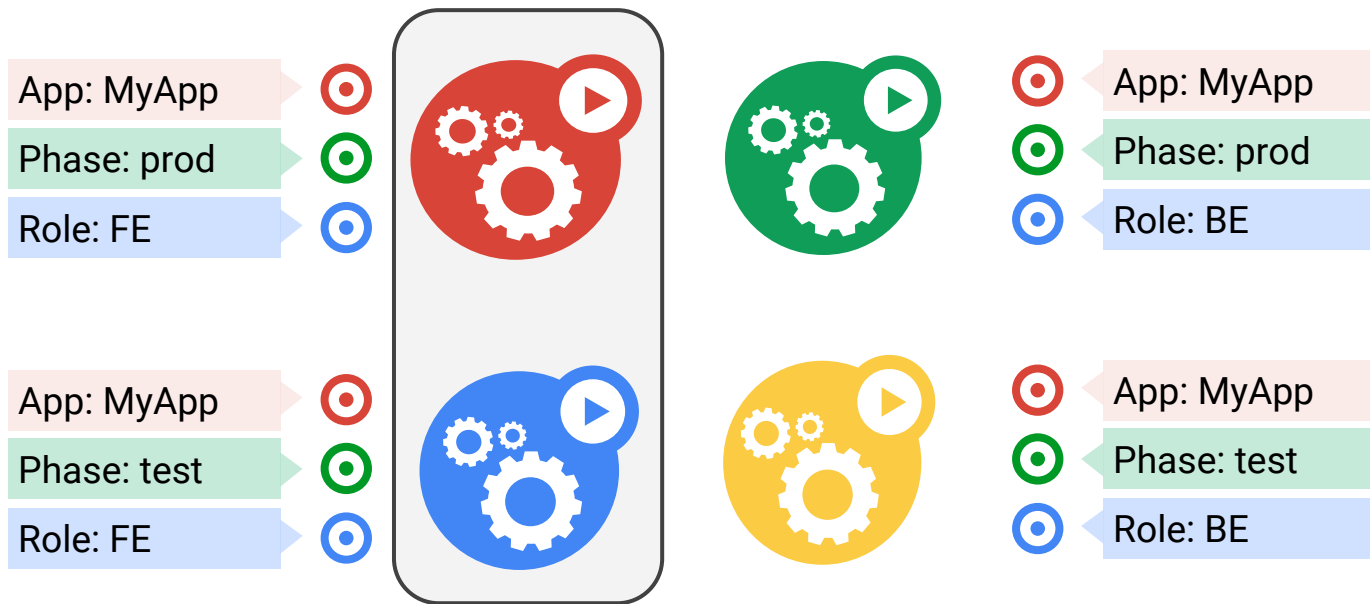


# You can query for labels that map to a value like the entire app



**App = MyApp**

Or narrow your search with multiple labels like your app's frontend



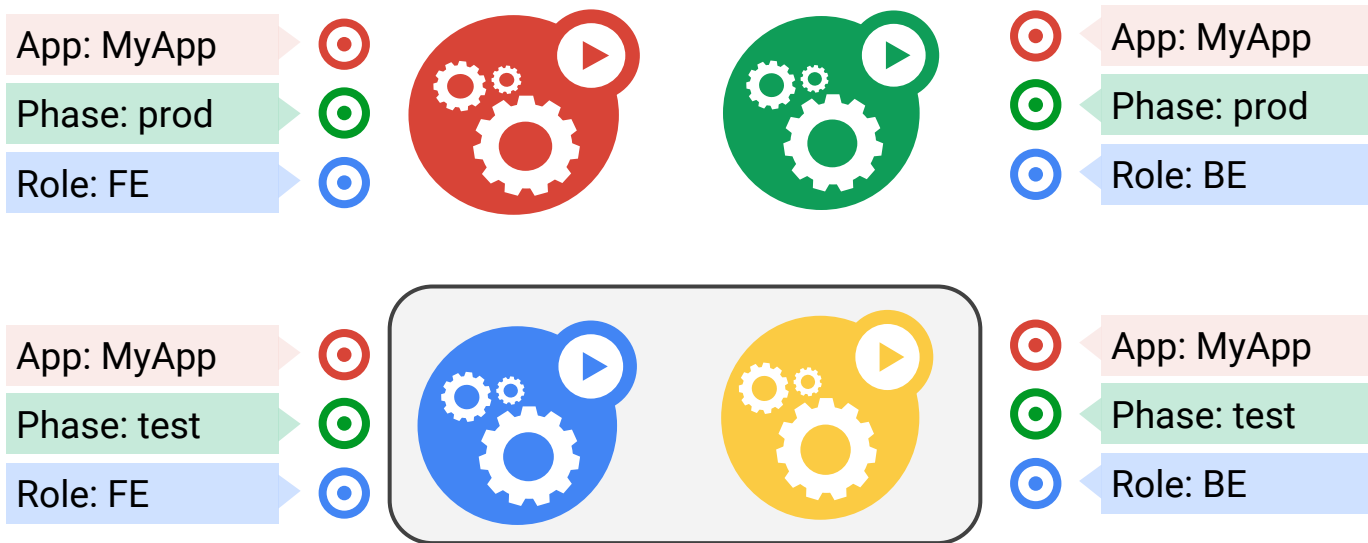
**App = MyApp, Role = FE**

## Or your app's backend



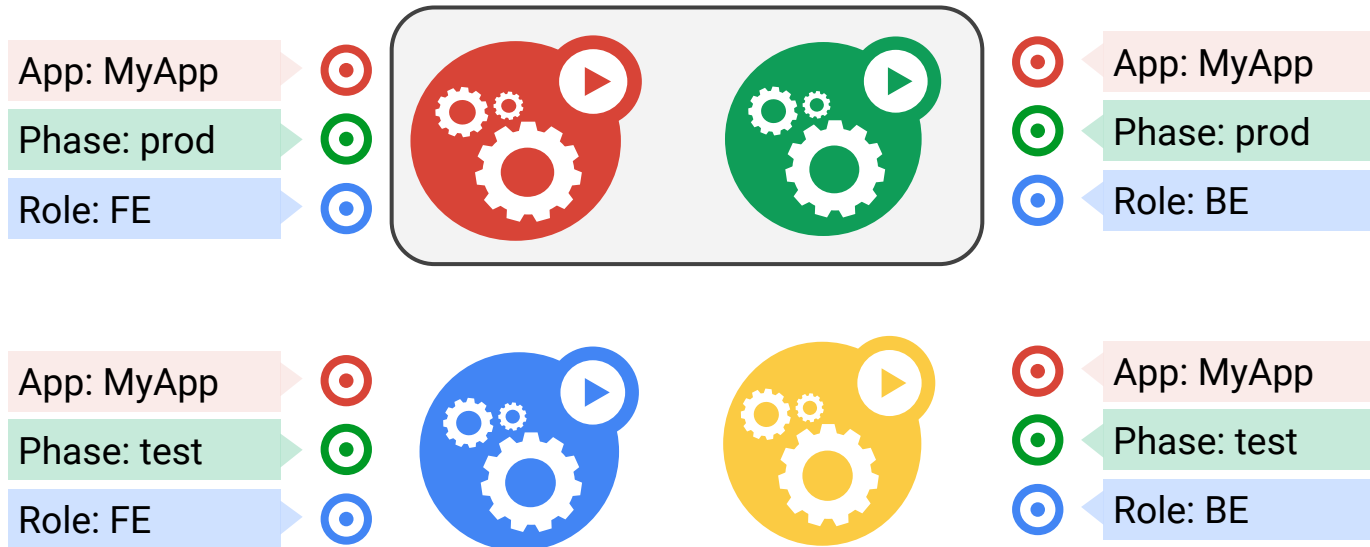
**App = MyApp, Role = BE**

# Or your app's test phase

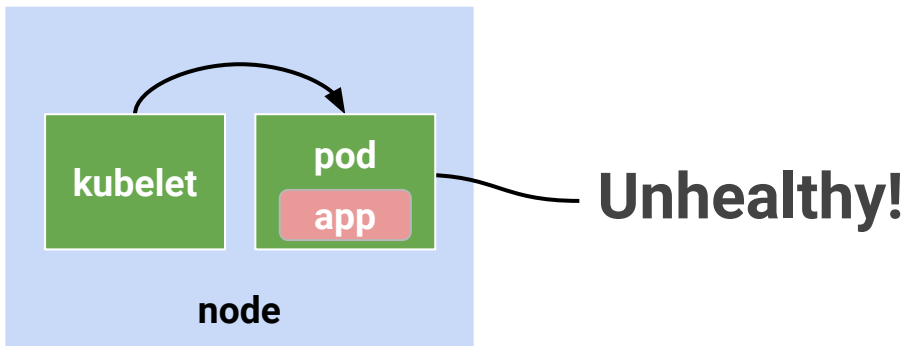




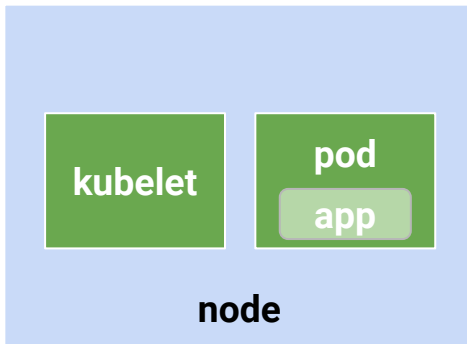
# Or your app's production release



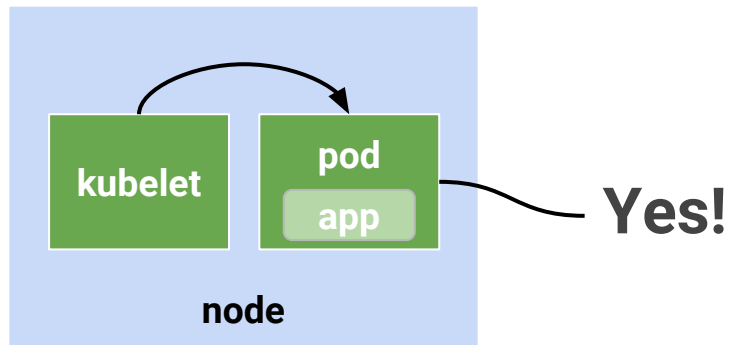
Kubelet checks whether the pod is alive and healthy; if it gets a negative response or no reply...



# Kubelet restarts the pod



# And continues until it gets a healthy reply



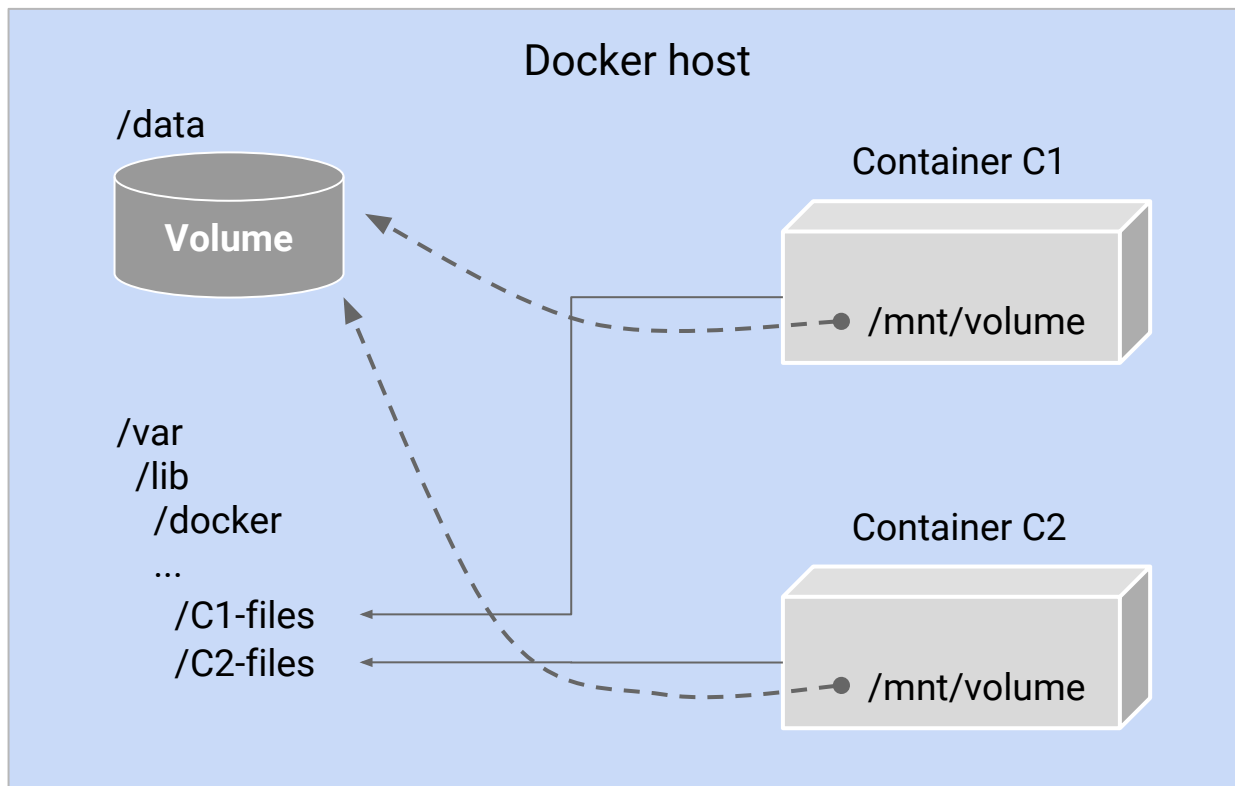
# Agenda

Clusters, nodes, and pods

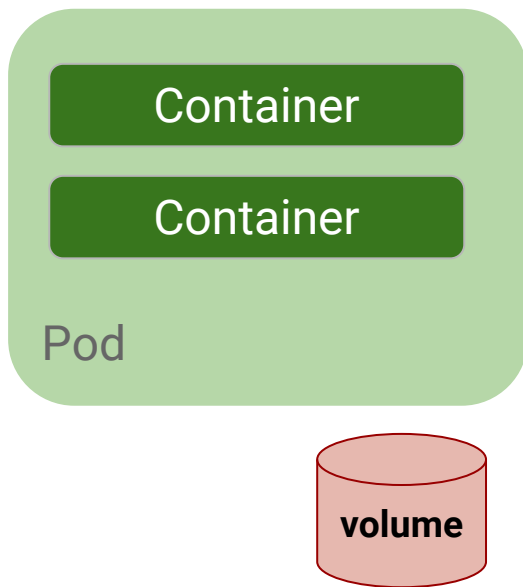
Services, labels, and selectors

Volumes

Docker provides data storage for containers, but volumes do not provide sharing between containers or lifecycle management

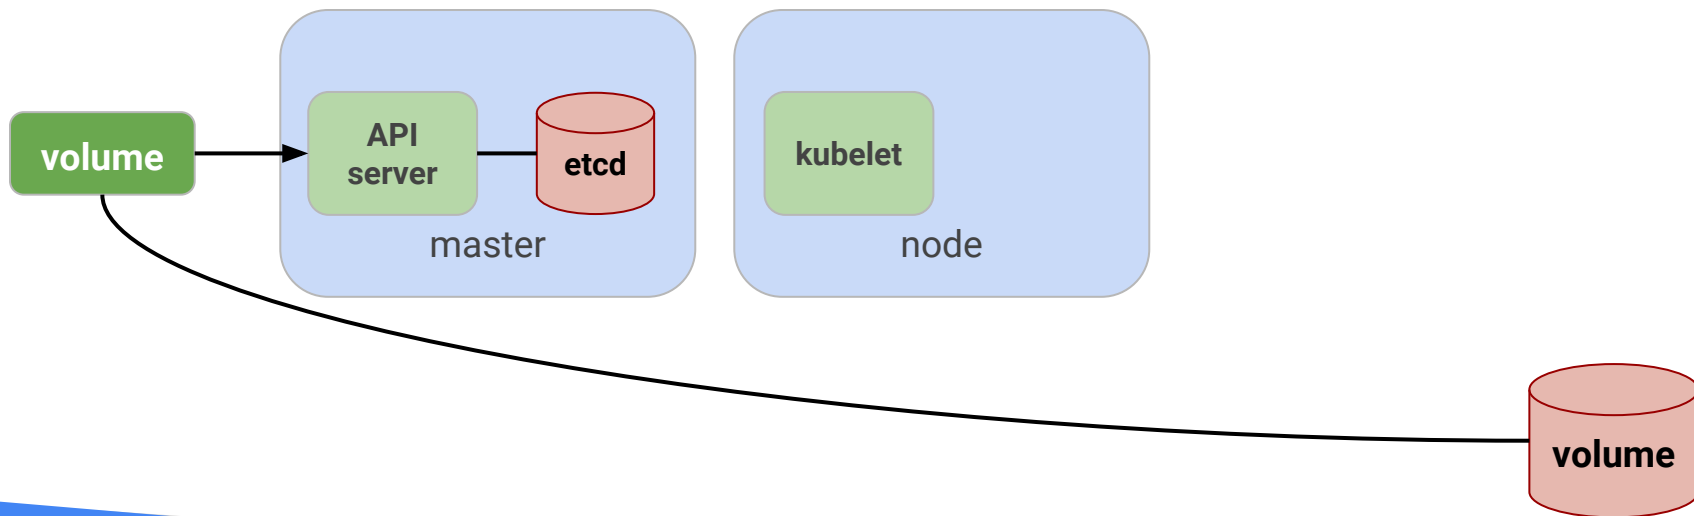


# Kubernetes volumes allow containers in pods to share data and be stateful



A volume is just a directory, and how it gets created depends on its type

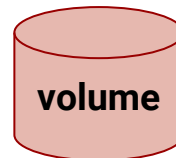
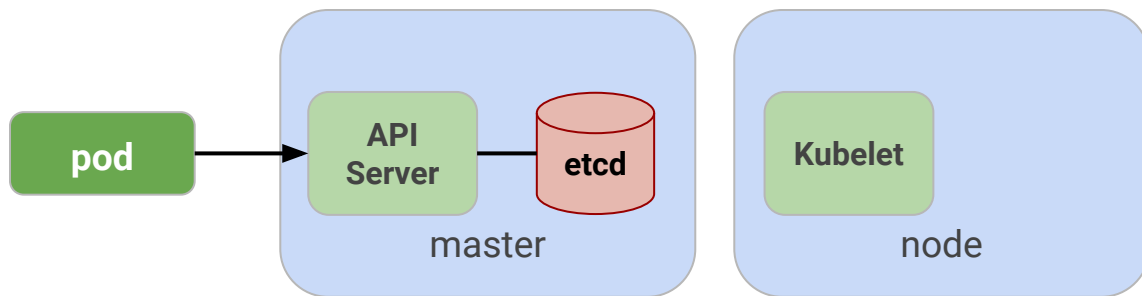
```
$> kubectl create <volume>
```



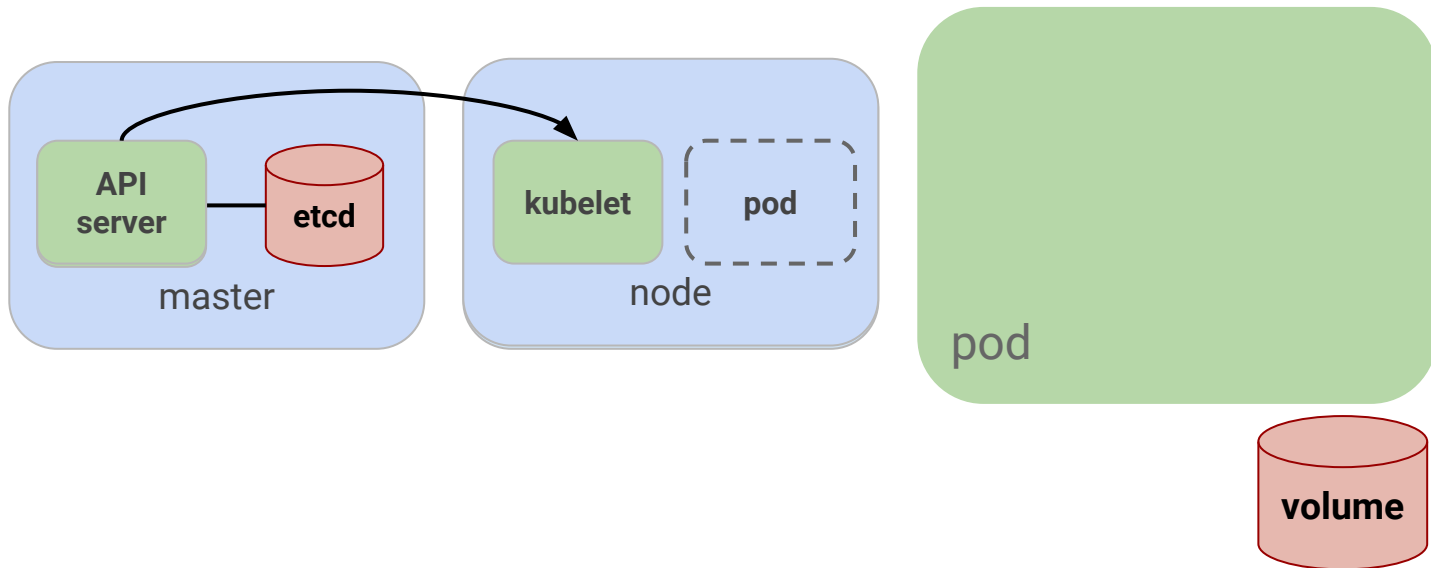


# Then you create a pod that consumes that data

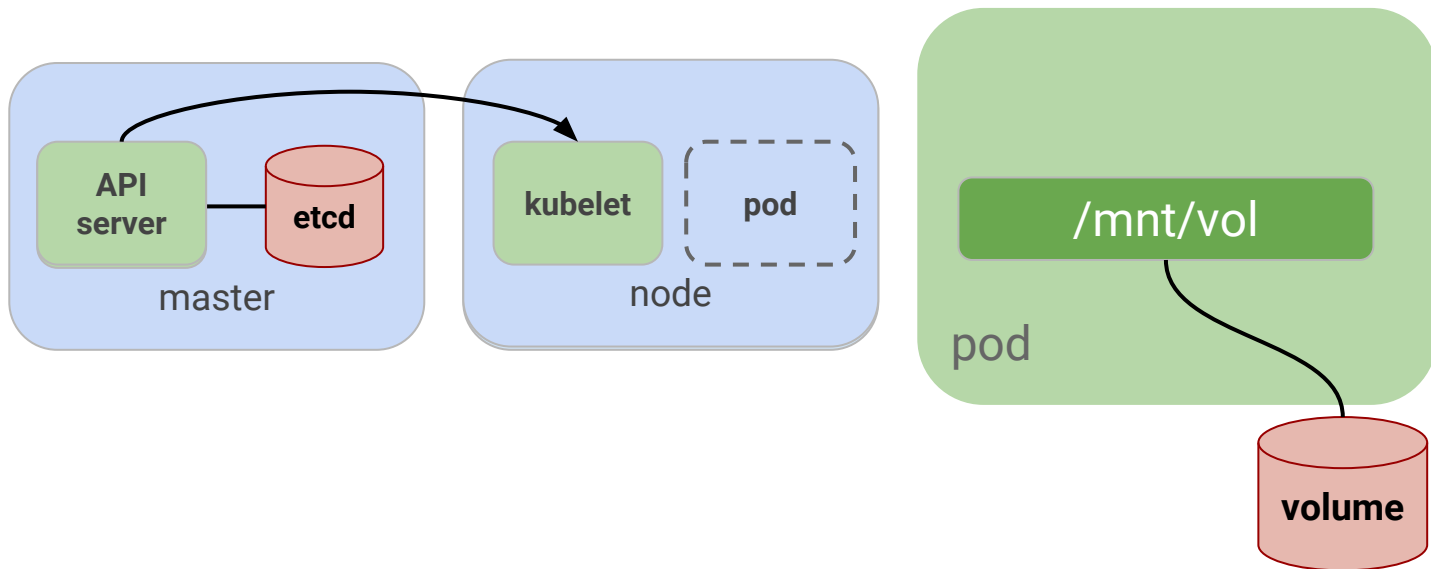
```
$> kubectl create -f pod.yaml
```



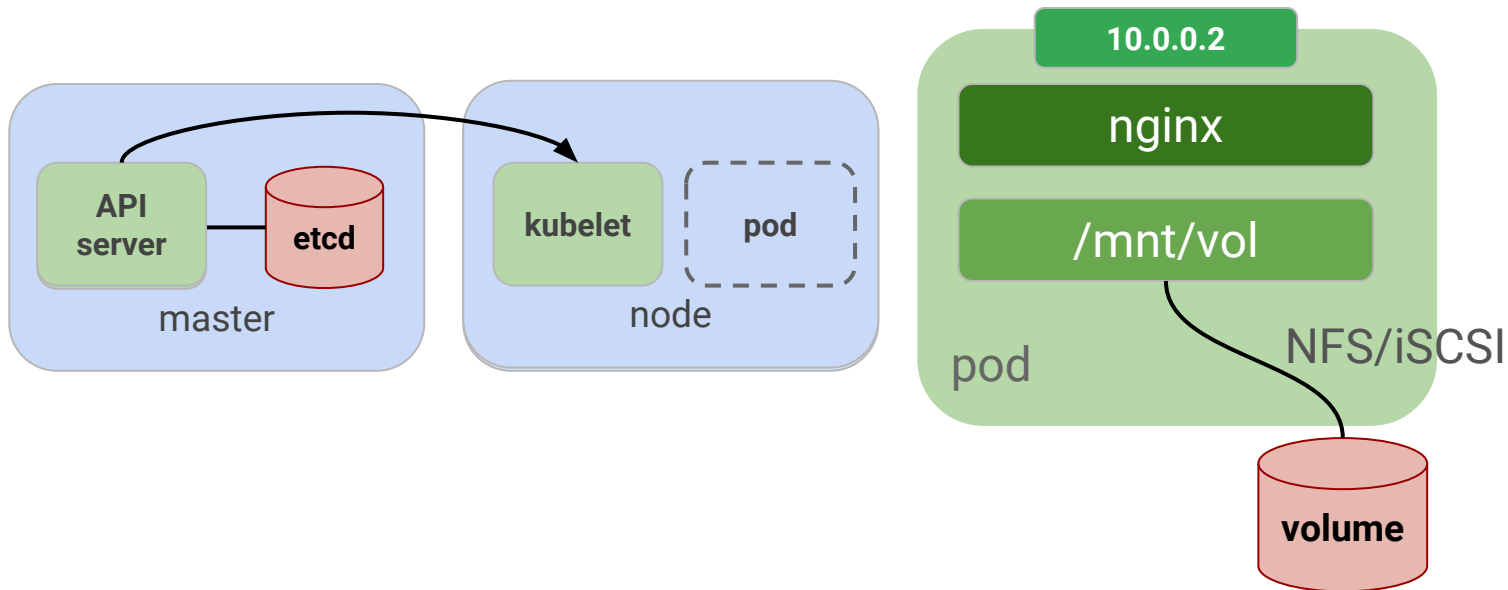
The volume is attached to the pod and made available to containers before they are brought online



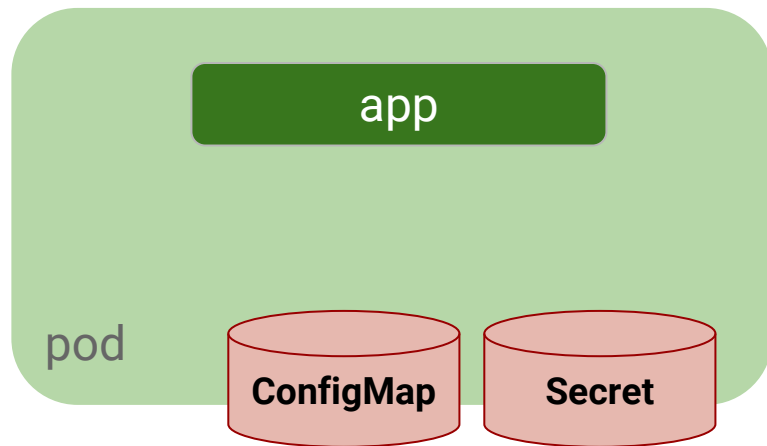
Once the volume is attached, data can be mounted into a container's file system



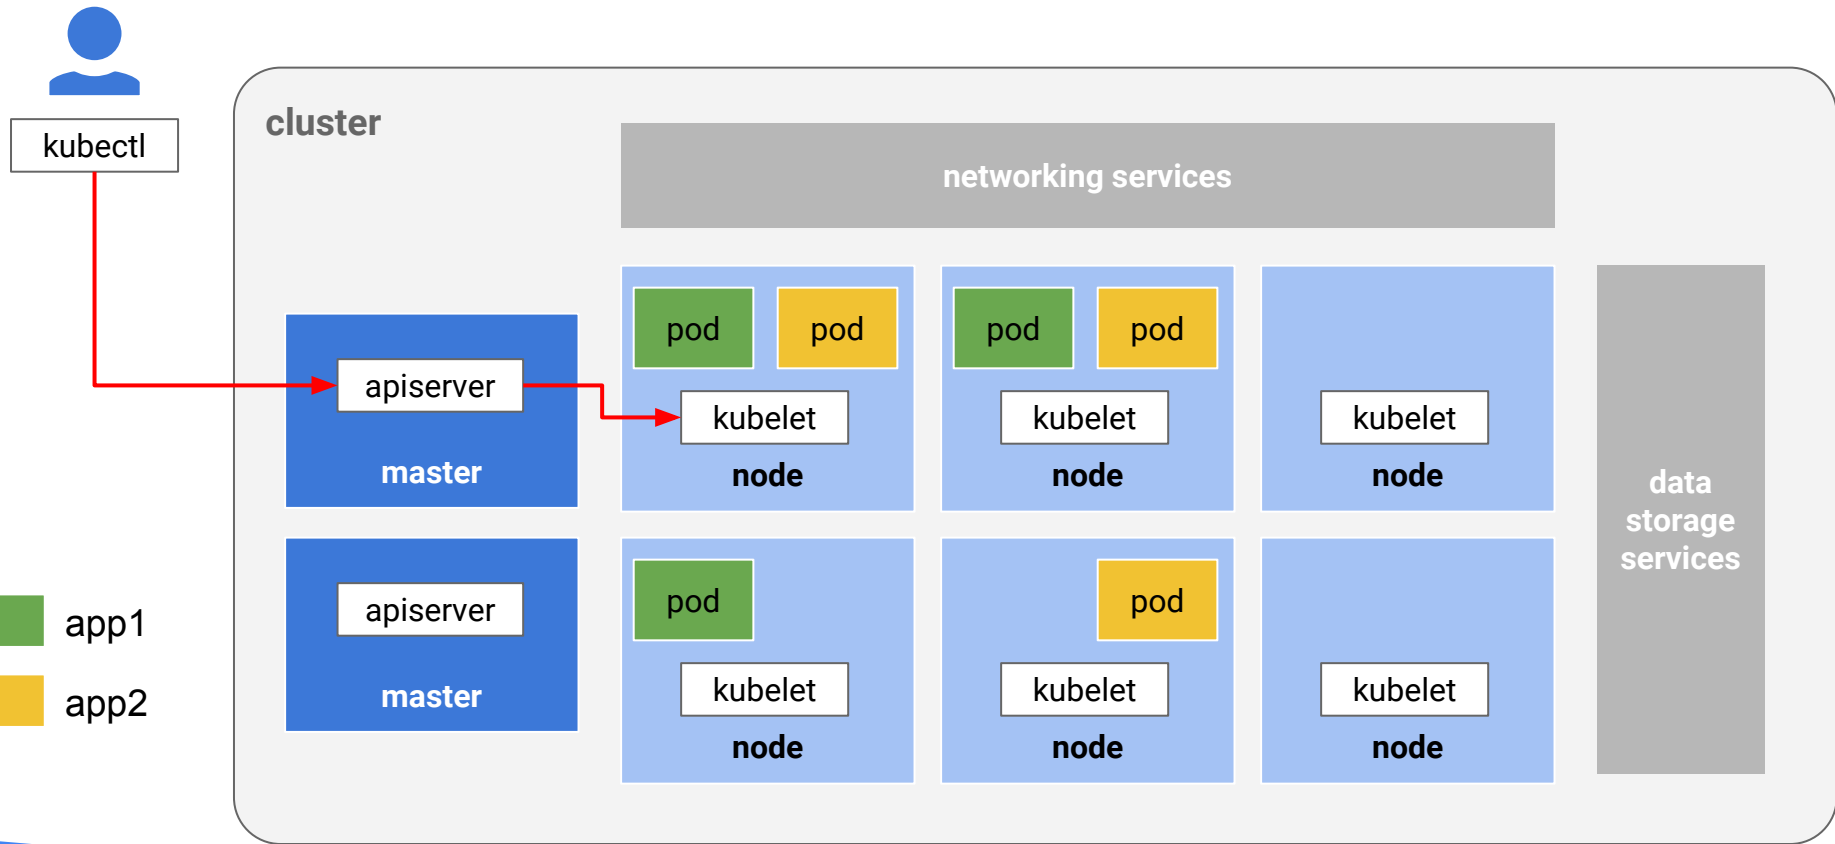
Then the container is run and can get the mounted data



# Some volumes share the lifecycle of their pod



# Here's a complete overview of a cluster



Lab